

# CSCI 551

# Numerical and Parallel

# Programming

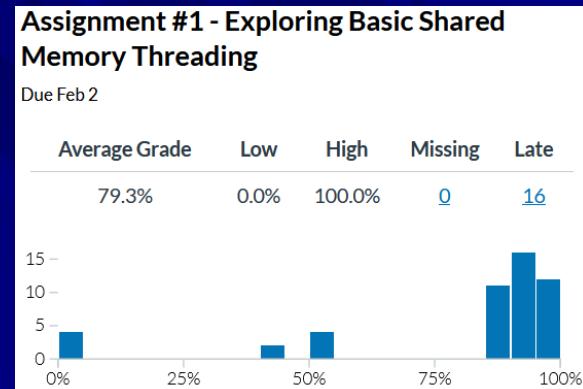
*Lecture 5-1 – Introduction to MPI and  
Simulation of Dynamic Systems*



# Assignment #1 & Quiz #1 – Results

## ■ Assignment #1

- Ave = 80.05% (a few no attempt made)
- Good progress and great start!
- Keep working on Amdahl's Law (P, SU, S)
- Methods to make parallel
- Analysis for speed-up
- Mr. G's law (saturates htop)
- Verifies (matches sequential output).
- Make sure you show results, analyze and explain them well, demonstrate understanding of numerical methods and code
- Average brought down by “no attempt” – start and turn in what you have done!



## ■ Quiz #1

- Quiz #1 **Don't forget to take before mid-night tonight if not already done!**
- Solutions posted after all have taken
- Questions on solutions? – Can cover on Thursday
- Should have basic understanding of Pthreads and OpenMP by now
  - Functional version of OpenMP as well as any well-structured block
  - When to use OpenMP and when to use Pthreads?
- Vector/Matrix and Prime Numbers – math covered so far

## ■ Grading policy - <csci551/policies/CSCI-551-Grading-Breakdown.pdf>

# Scalable Computing Resources

## ■ NSF ACCESS – Largest scale possible

The screenshot shows the NSF ACCESS portal interface. It displays two active projects:

- CIS240363: Chico Research for Environmental Science and Technology: Elephant Listening Anti-Poaching sensor fusion project**:
  - Explore: May 29, 2024 to May 28, 2025
  - Overview tab selected.
  - Resource table:
    - SDSC Expanse GPU (Active)
    - SDSC Expanse Projects Storage (Active)
  - Role table:
    - PI: Samuel Siewert
- CIS250125: Quantum vs. Parallel Advantage - A Study of Semi-Prime Factoring with CUDA and CUDA-Q**:
  - Explore: Feb 13, 2025 to Feb 12, 2026
  - Overview tab selected.
  - Resource table:
    - MATCHPlus (Active)
  - Action Details table:
    - New: Feb 12, 2025 (Approved)

A large black box highlights the CIS250125 project area, with the text "Parallel + Quantum" overlaid on it.

## ■ CSU Chico:

- NUC Cluster: 62 nodes - 136 SMT cores
- ARM Clusters: Jetson Nano cluster (20 nodes, 80 ARM cores, 2,560 SP co-procs), R-Pi-4 cluster (40 nodes, 160 cores)
- GP-GPU: A100 80GB, 512GB with 16 Xeon cores
- ESYS: DGX-1 V100 Infiniband cluster (4 nodes).

# Exercise #3 – Features MPI

- MPI introduced with Exercise #3
- Read Pacheco, Chapter 3.1 – 3.3 (MPI introduction)
- See MPI examples:
  - [csci551/code/hello\\_cluster/](#)
  - [csci551/code/MPI\\_Examples/](#)
- Make sure you understand ECC-Cluster and CSCI
  - [csci551/README-cluster.html](#)
  - See notes on CSCI cluster (Jetson Nano and R-Pi 4)
  - Test hello\_cluster on both!
  - Get help from IT or during office hours if you have account issues!

# CSCI Cluster and ECC-Cluster Update

- Appears to be working with firewall work-around
- Still testing ... works with our starter code with mpiexec
- mpiexec should always work
- mpirun is a bit more complex ...

```
sbsiewert@rpil:~/csci551/numeric-parallel-starter-code/hello_cluster$ mpirun --mca btl_tcp_if_include eth0 -n 8 -host rpil:4 -ho
st rpi3:4 ./greetings
Hello from process 0 of 8 on rpilHello from process 1 of 8 on rpil
Hello from process 2 of 8 on rpil
Hello from process 3 of 8 on rpil
Hello from process 4 of 8 on rpi3
Hello from process 5 of 8 on rpi3
Hello from process 6 of 8 on rpi3
Hello from process 7 of 8 on rpi3
sbsiewert@rpil:~/csci551/numeric-parallel-starter-code/hello_cluster$ █
```

- Running with mpirun on ECC-Linux simpler due to difference in firewall and set-up for SSH perhaps

# Tips: Don't Suffer in Silence & Late Work

- Use 3-days grace when life hits-you-over-the-head – do not abuse it (10% penalty for abuse)
- Come to office hours ([Sam-Siewert-Fall-2023-Office-Hours.pdf](#)) or make an appointment
  - Don't rat-hole, e.g., "have over 20 hours sunk into this ...", this is a mistake to not seek out proper help!
  - Get pointed in the right direction
- Asking for help 2 days after assignment due date is not early enough – review assignment when posted, then chip away at it
- We have 2 weeks per assignment so plenty of time to seek out help before the due date including clarifications
- Suffering in silence in industry and noting issues after due dates can get you in trouble. Better to learn now!

# Summary of Image Processing - Scaling

- Work from 2021 by 551 student, ERAU students and myself, submitted to IEEE Aerospace, presentation at Big Sky in March
- Scaling with Power Efficiency (Pthreads, OpenMP, CUDA)
- For Ground, Aerial, and Space Robotics

<https://github.com/sbsiewertcsu/CISA-software-benchmarks>

## Common Instrument Stack Architecture Study



January 31, 2021



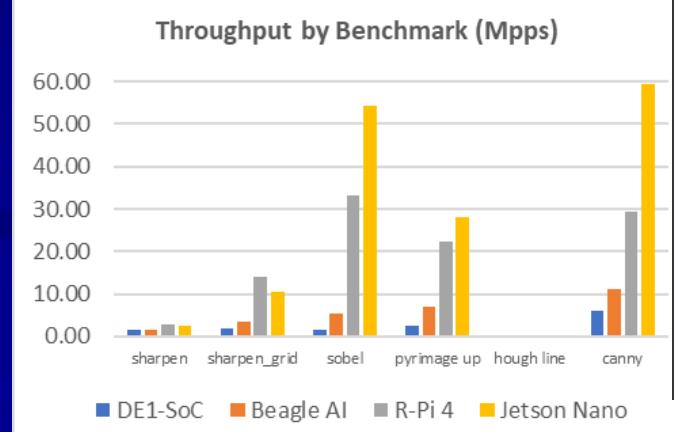
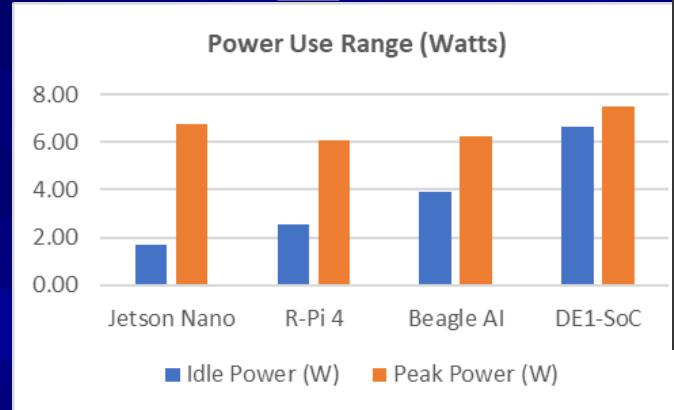
© Sam Siewert, ICARUS & EMVIA Group  
IEEE Aerospace Conference, Big Sky, 2021

## IEEE Aerospace

Student Research Group on THURSDAYS @ 6:30PM PDT, with ZOOM

### Power Use:

1. IoT
2. Sensor Networks
3. Edge AI and CV



### Pthreads, OpenMP & CUDA

1. Saturate CPU cores
2. Saturate Co-Pros

# Terminology Discussion

## ■ What is the difference between?

- Process (definitely a Container – address space)
- Thread (at least a Sequence of instructions + what?)
- Task (in case you are interested)

## ■ Think about examples, address space, processor nodes in clusters, execution context, context switching, sharing data, and sharing code

## ■ Optional Activity: Break-out and discuss

## ■ Report Back: Summarize each

### Process:

- Protects apps from rogue apps
  - Bad pointers
  - Stack overflow
  - Non-recoverable exceptions (e.g.,  $x=n/0;$ )
- Contains at least 1 thread
- Can contain many threads
- Keeps track of resources
  - Open files “lsof”
  - OS mechanisms – e.g., message queue
  - Memory use (heap, stack, data, code, bss)
  - Child processes created (fork, exec)
  - Standard input, output, error
  - Error numbers and state
  - Timers and timeouts
  - Blocked process/thread
  - Signals

# Working Definitions

- **Process** – an address space “container” for a thread of execution with stack, register state, and PC state along with significant additional software state such as copies of all I/O descriptors (much more than a task control block for example) **including a protected memory data segment** (protected from writes by other processes). A process may be implemented by a kernel level task (in Linux).
- **Thread (of execution)** - a thread is the execution trace and context of a CPU’s PC (Program Counter) over time not including *context switch* code execution by an RTOS or OS. State information may or may not be associated with a thread of execution, but the value of the PC before a context switch is the minimum state that must be maintained on a system which includes *preemption*, *including simple interrupts*.
  - A thread might be as simple as a boot thread, running code out of ROM (Read Only Memory), with no stack, data segment, or anything other than a sequence of machine code instructions.
  - More fully featured user space threads, like Pthreads, have stack, data segments, register state (saved and restored on context switches) and are a method of providing concurrent features without the overhead of a full process in user space and **Pthreads share the address space of a process, and are contained within a process**.
- **Task** – an “execution context” for a thread with normal thread state including stack, registers, PC (Program Counter), but also including signal handlers, task variables, task ID and name, priority, entry point, and a number of state and inter-task communication data contained in a TCB (Task Control Block). Tasks are the most common execution context used in an RTOS and are also a common OS kernel execution context.

# MPI – Message Passing Interface

## ■ Software Scaling and Parallel Computation with Clusters

- Clusters are nodes interconnected via a network
- Network may range from fully connected to mesh to ring
- Most modern systems interconnected with a mesh
- Either a Torus (e.g., Infiniband) or a non-blocking switch (COTS cluster)
- Non-blocking, line-rate switches for low port count, low data rate, readily available
- Non-blocking, line-rate switches for high port count, high data rate, are expensive \$\$\$



16 port, \$1,100 on Amazon – 2020  
< \$100 per port, **10G uplink**  
**Gigabit < \$10 oer port**

## MPI (Message Passing Interface)

Messages point to point via switch (without going off switch)

Outside connection is simply to get on the cluster

May need to upload/download data sets

# MPI – SPMD (pp. 83-113 Pacheco)

- The basic concept is a Single Program that creates multiple data processing processes
  - Processes are created on one node and/or distributed over all nodes in the cluster
  - Messages provide data to each process
  - Each process has a unique identification
  - Each can then work on a portion of the parallel problem
- Results can be merged at the end
- Reduction operators – e.g., summing, maximum, can be applied and create a barrier for ranks

# Before Attempting MPI Coding

- Read Instructions - [csci551/README-cluster.html](#)
- Check path to required Intel compiler tools
  - /opt/intel/compilers\_and\_libraries\_2020.0.166/linux/mpi/
  - Compiler “mpicc” – intel64/bin/mpicc
  - Cluster tool “mpirun” – /opt/intel/intelpython3/bin/mpirun
  - Single node tool “mpiexec” – /opt/intel/intelpython3/bin/mpiexec
  - C++ Compiler “mpicxx” - intel64/bin/mpicxx
- Make sure you can SSH to each node used without password prompt
  - Otherwise mpirun will just HANG
  - Machines being used should all be reachable too
    - [csci551/c1 hosts test.sh](#)
    - [csci551/c2 hosts test.sh](#)

# MPI Demonstrations (CSU & Intel)

## ■ MPI greetings with defaults

```
sbsiewert@o244-01:~/code/hello_custer$ ls  
cl_machine_file c2_machine_file greetings greetings.c Makefile  
sbsiewert@o244-01:~/code/hello_custer$ mpirun -n 4 ./greetings  
Hello from process 0 of 4 on o244-01  
Hello from process 1 of 4 on o244-01  
  
Hello from process 2 of 4 on o244-01  
  
Hello from process 3 of 4 on o244-01  
sbsiewert@o244-01:~/code/hello_custer$ █
```

```
sbsiewert@o244-01:~/mpi-benchmarks-master$ mpirun -n 16 ./IMB-MPII  
#-----  
#   Intel(R) MPI Benchmarks 2019 Update 6, MPI-1 part  
#-----  
# Date : Tue Sep 22 16:33:58 2020  
# Machine : x86_64  
# System : Linux  
# Release : 4.15.0-117-generic  
# Version : #118-Ubuntu SMP Fri Sep 4 20:02:41 UTC 2020  
# MPI Version : 3.1  
# MPI Thread Environment:  
  
# Calling sequence was:  
# ./IMB-MPII  
  
# Minimum message length in bytes: 0  
# Maximum message length in bytes: 4194304  
#  
# MPI_Datatype : MPI_BYTE  
# MPI_Datatype for reductions : MPI_FLOAT  
# MPI_Op : MPI_SUM  
#  
#  
# List of Benchmarks to run:  
# PingPong  
# PingPing  
# Sendrecv  
# Exchange  
# Allreduce  
# Reduce  
# Reduce_local  
# Reduce_scatter
```

Intel Makefile Updated  
for CSU cluster nodes

<https://www.ecst.csuchico.edu/~sbsiewert/csci551/code/mpi-benchmarks-master-csu.zip>

# OpenMP Demonstrations (Intel CSU)

- [csci551/ipsxe 2019 html/compiler\\_c/openmp\\_samples/readme.html](https://csci551.ipsxe.2019.html/compiler_c/openmp_samples/readme.html)
- [csci551/code/openmp\\_samples/intel\\_csu.zip](https://csci551.code/openmp_samples/intel_csu.zip)

```
sbsiewert@csci551$ cd openmp_samples_intel_csu/
sbsiewert@openmp_samples_intel_csu$ ls
DCT  Mandelbrot  Matrix_Multiply  MergeSort  readme.html
sbsiewert@openmp_samples_intel_csu$ cd DCT
sbsiewert@DCT$ make run
icpc -c -g -O2 -xAVX -fopenmp -std=c++11 -o release/matrix.o src/matrix.cpp
icpc -c -g -O2 -xAVX -fopenmp -std=c++11 -o release/timer.o src/timer.cpp
icpc -c -g -O2 -xAVX -fopenmp -std=c++11 -o release/DCT.o src/DCT.cpp
TBB Warning: tbb.h contains deprecated functionality. For details, please see Deprecated Features appendix in the TBB reference manual.
Linking...
icpc release/matrix.o release/timer.o release/DCT.o -fopenmp -tbb -o release/DCT
./release/DCT res/nahelam.bmp res/nahelam1.bmp
Please enter the version you want to execute:
0) All versions
1) Scalar version without parallelism
2) OpenMP SIMD without parallelism
3) Scalar version with TBB parallelism
4) OpenMP SIMD with TBB parallelism
4
The time taken in number of ticks is 354533566
The time taken is 0.222091 seconds
sbsiewert@DCT$
```

## Intel Threading and OpenMP DCT

```
sbsiewert@Matrix_Multiply$ make run
icc -c -O2 -std=c99 -wd3180 -o out/openmp_sample.o src/openmp_sample.c
icc out/openmp_sample.o -o out/openmp_sample
Problem size: c(600,2400) = a(600,1200) * b(1200,2400)
Calculating product 20 time(s)
```

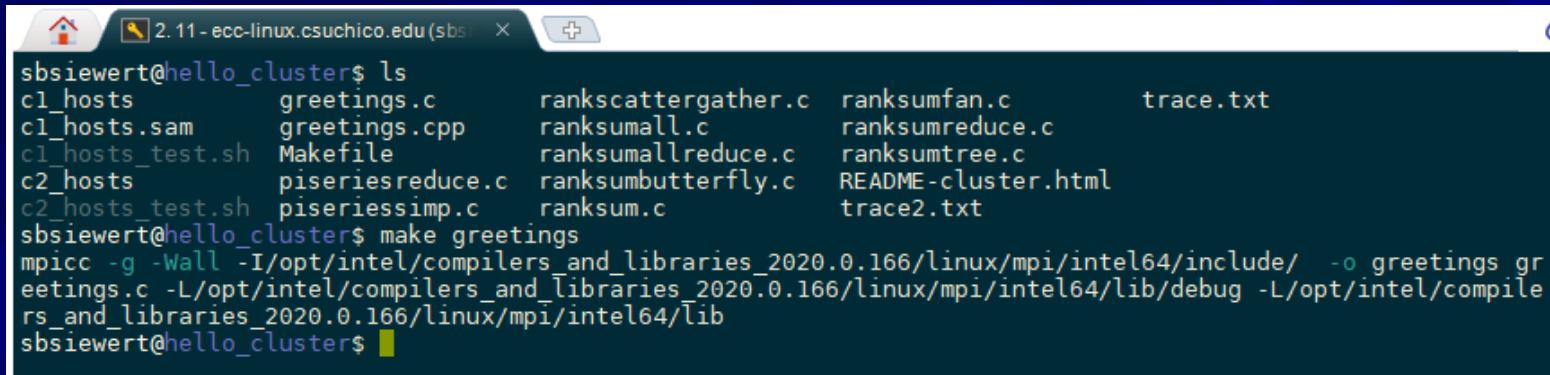
We are using 1 thread(s)

```
Finished calculations.
Matmul kernel wall clock time = 20.21 sec
Wall clock time/thread = 20.21 sec
MFlops = 3419.477911
sbsiewert@Matrix_Multiply$
```

## Intel OpenMP Matrix Multiply

# Simple MPI Examples – Build and Run

- Make like any other program, but using mpicc and note special includes (-I) and libraries (-L)



A screenshot of a terminal window titled "2.11-ecc-linux.csuchico.edu (sbs...)" showing the build process for MPI examples. The user runs "ls" to list files, then "make greetings" to compile the "greetings.c" file. The command "mpicc -g -Wall -I/opt/intel/compilers\_and\_libraries\_2020.0.166/linux/mpi/intel64/include/ -o greetings greetings.c -L/opt/intel/compilers\_and\_libraries\_2020.0.166/linux/mpi/intel64/lib/debug -L/opt/intel/compilers\_and\_libraries\_2020.0.166/linux/mpi/intel64/lib" is shown in green, indicating it's being typed or has been run.

```
sbsiewert@hello_cluster$ ls
cl_hosts      greetings.c      rankscattergather.c   ranksumfan.c      trace.txt
cl_hosts.sam  greetings.cpp    ranksumall.c        ranksumreduce.c
cl_hosts_test.sh  Makefile     ranksumallreduce.c  ranksumtree.c
c2_hosts      piseriesreduce.c ranksumbutterfly.c README-cluster.html
c2_hosts_test.sh  piseriessimp.c ranksum.c          trace2.txt
sbsiewert@hello_cluster$ make greetings
mpicc -g -Wall -I/opt/intel/compilers_and_libraries_2020.0.166/linux/mpi/intel64/include/ -o greetings greetings.c -L/opt/intel/compilers_and_libraries_2020.0.166/linux/mpi/intel64/lib/debug -L/opt/intel/compilers_and_libraries_2020.0.166/linux/mpi/intel64/lib
sbsiewert@hello_cluster$
```

- Run with mpirun and specify nodes to use with “-f”

```
sbsiewert@hello_cluster$ mpirun -n 8 -ppn 2 -f cl_hosts ./greetings
Hello from process 0 of 8 on o251-01
Hello from process 1 of 8 on o251-01
Hello from process 2 of 8 on o251-02
Hello from process 3 of 8 on o251-02
Hello from process 4 of 8 on o251-03
Hello from process 5 of 8 on o251-03
Hello from process 6 of 8 on o251-04
Hello from process 7 of 8 on o251-04
sbsiewert@hello_cluster$
```

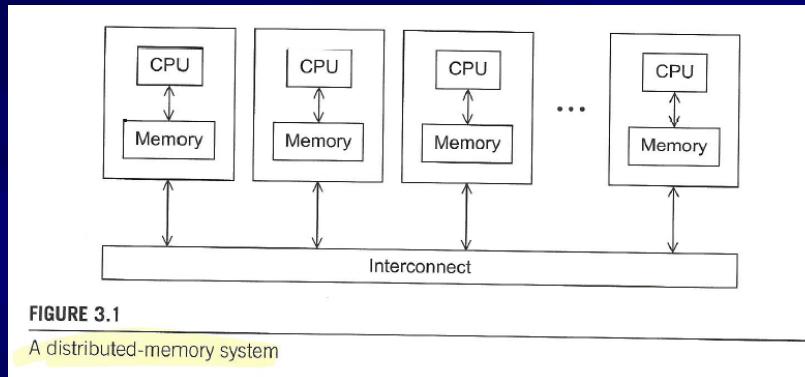
Follow Steps in:  
[csci551/README-cluster.html](#)

Test node status:  
[csci551/c1\\_hosts\\_test.sh](#)  
[csci551/c2\\_hosts\\_test.sh](#)

Use: [c1\\_hosts](#) & [c2\\_hosts](#)

# Recall Shared vs. Distributed Memory

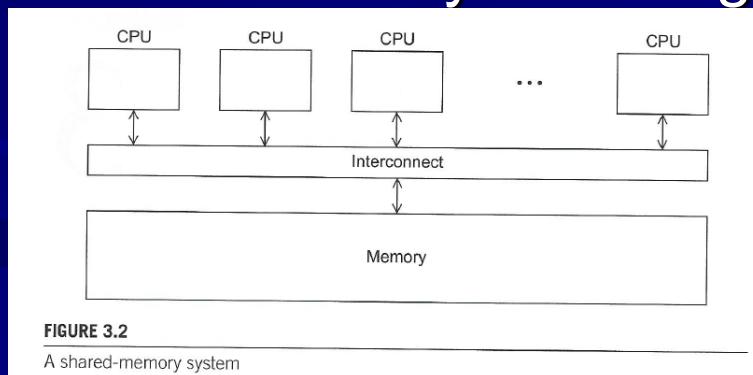
- Each process is a unique address space
- Each node has a unique memory device



Interconnect – e.g., switch

CPU + Memory – cluster node

- Shared memory on a single node



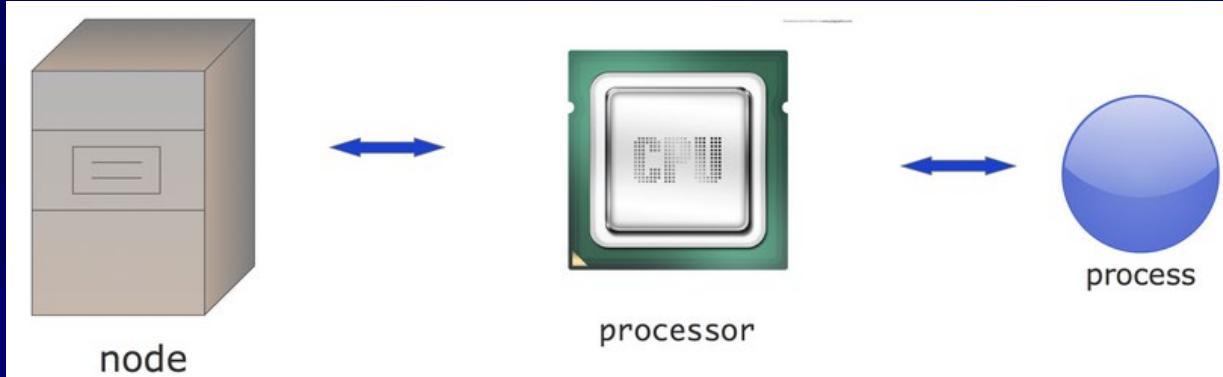
Interconnect

- Memory bus
- Scaling bus (e.g., QPI)

Core + Cache – CPU core  
CPU + QPI – CPU socket

# Recall Scale-up – Add CPUs with cores

## ■ Scaling up a single node

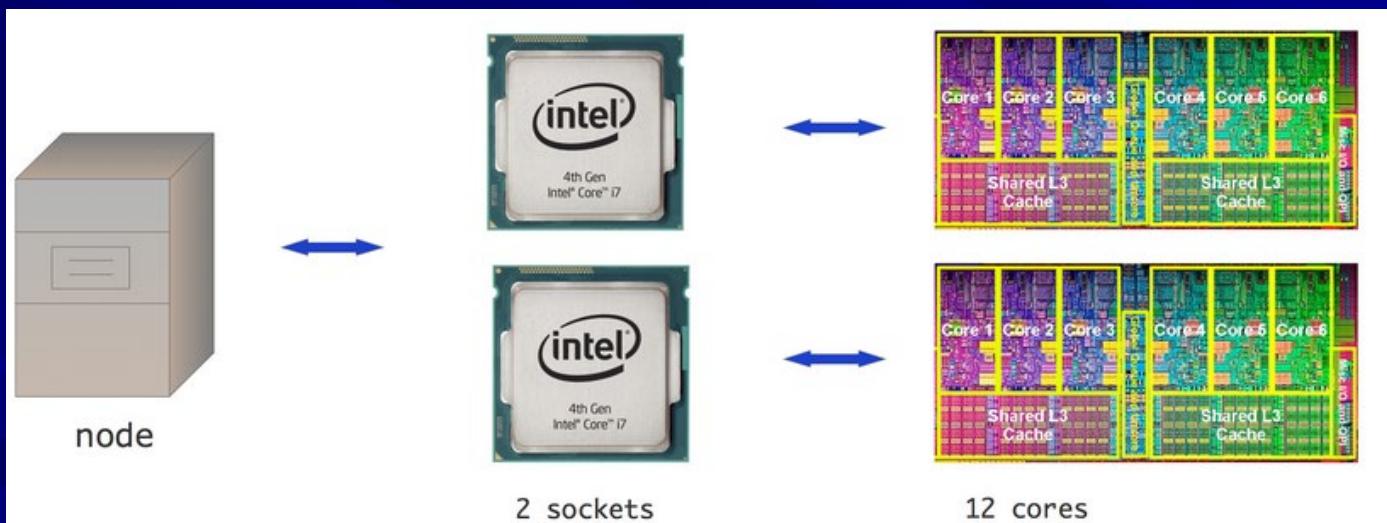


Cost per Core

Intel i7 - \$60/core, \$30 SMT

ARM R-Pi - \$20/core

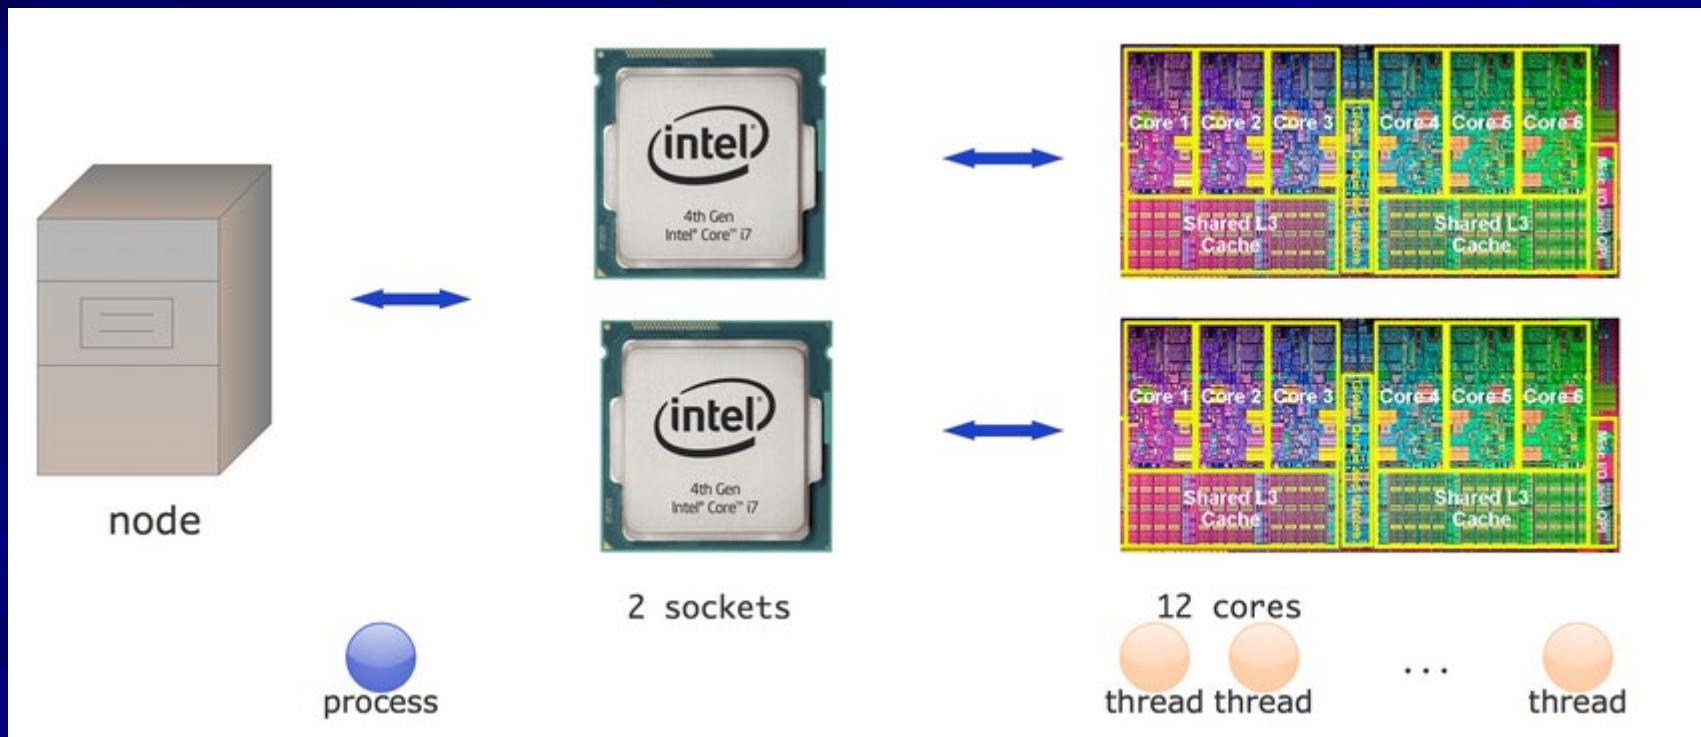
ARM Jetson - \$20/core, \$1/Co-Proc



<https://pages.tacc.utexas.edu/~eijkhout/pcse/html/mpi-functional.html>

# Recall Scale-up – Shared Memory Pthreading

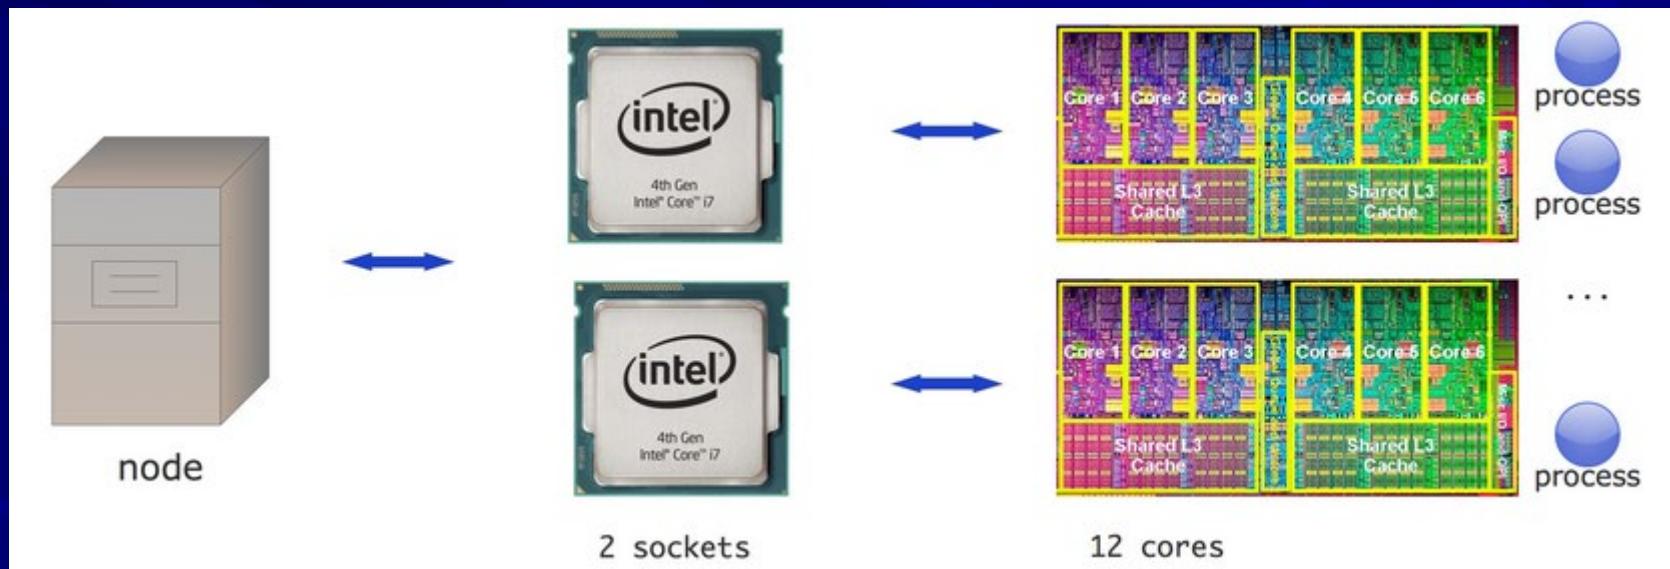
- Scaling up a single node and multi-threading process
  - Shared address space within a single process
  - Threads run on multiple CPUs and processor cores in one address space



<https://pages.tacc.utexas.edu/~eijkhout/pcse/html/mpi-functional.html>

# Recall Scale-up – Multi-programming

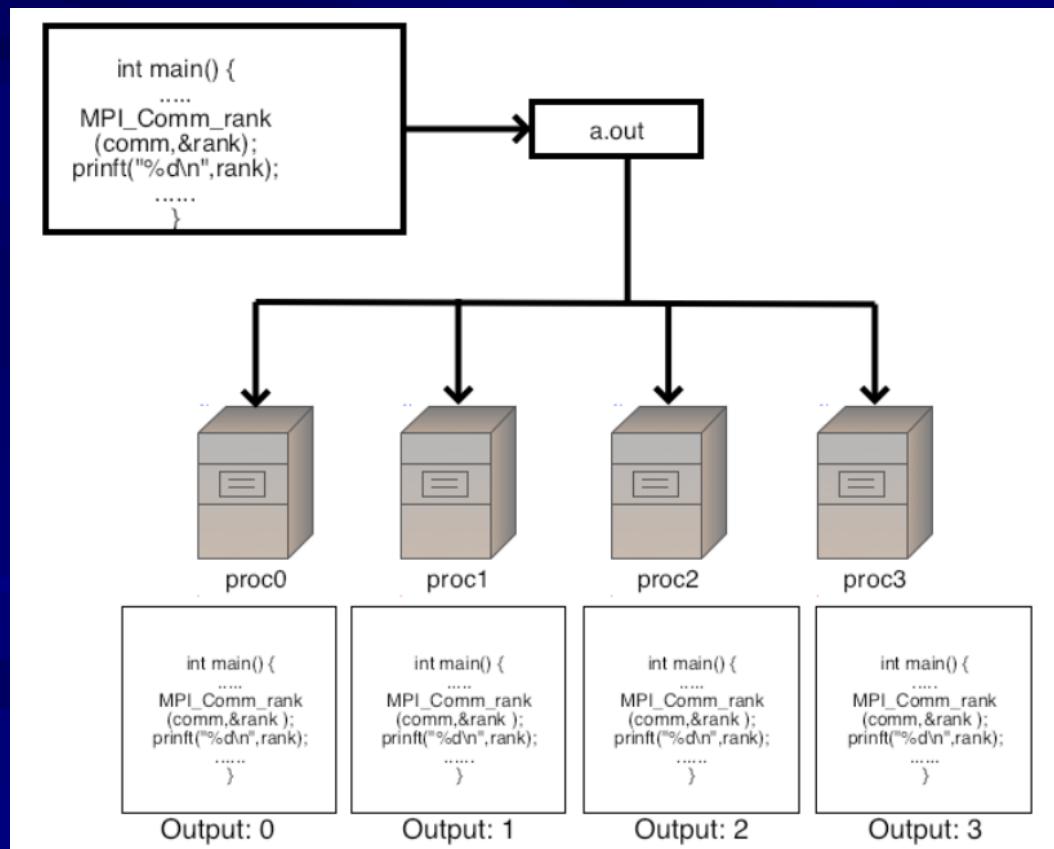
- Scaling up a single node and load balancing processes
  - Separate address space for each process (use messages to communicate and synchronize)
  - Processes run on multiple CPUs and processor cores



<https://pages.tacc.utexas.edu/~eijkhout/pcse/html/mpi-functional.html>

# Scale-out to Multiple Nodes – ECC Cluster

- As I create more processes to divide up work, assign each process to Node, CPU, and core
- Use MPI to divide up work, identify role, merge results



## Linux nodes

Each node schedules processes and provides load balancing

## SMP UMA/NUMA nodes

Details of node selection and CPU/core selection are provided by system

## Cluster + OS

# Comparison Discussion

- Which provides most cost-effective and efficient scaling?
  - Distributed memory, cluster scaling with MPI
  - Shared memory, CPU and core scaling with Pthreads or OpenMP
  - Hybrid using both?
- Which is simpler to program?
- Which is easier to use for speed-up?
- Which can be scaled the most?
  - In terms of number of processors and processes/threads

# Exercise #3 – Get Started Early!

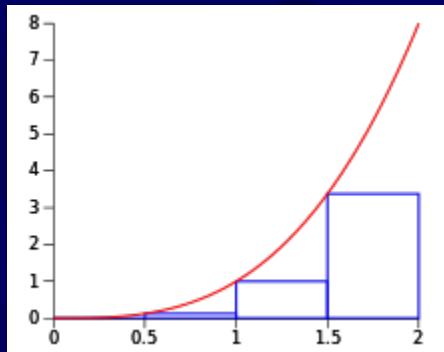
- **Read chapter 3 through p. 113 – Intro to MPI**
- **Starter - [csci551/code/functiongen/](#)**
- **Pacheco Code - [csci551/code/MPI Examples/trap.c](#)**
- Review [LLNL tutorials](#) (for more on Pthreads as you are interested)
- Review Intel [MPI benchmark user's guide](#)
- Simple MPI program – based on book to explore MPI and cluster features
- Top Down - MPI benchmarks
  - Build, run, observer
  - Review code and explain
- Bottom Up – Methods of integration
  - Simple forward integration, large, medium, small step size
  - Trapezoidal integration, large, medium, small step size
- Application – Train Acceleration and Braking
  - Create a profile of position and velocity over time
  - Given acceleration timeline
  - Compute arrival time
  - Use MPI to speed up the analysis (6 runs, “n” segments to integrate)

# Exercise #3 – MPI for Simulation

- Continuous physics of 1-DOF translation
  - Position =  $P_0 + \text{Velocity} \times \text{time}$ , if  $V(t)$  is constant
  - Velocity =  $V_0 + \text{Acceleration} \times \text{time}$ , if  $A(t)$  is constant
  - Acceleration can be constant, or could be a function of time
- If Acceleration is  $f(t)$ , we need to integrate
  - Simple Riemann Sum Integration
  - $\text{Position}(t) = P_0 + \sum V(t)dt$
  - $\text{Velocity}(t) = V_0 + \sum A(t)dt$
  - $\text{Acceleration}(t) = f(t)$ , provided as acceleration profile (or from sensor)
- Approximate the area under curve  $A(t)$  for a  $dt$  (step size)
- Break  $A(t)$  into segments or make it a look-up-table
- Interpolate for  $A(t)$  at times between profile given
- Consider accuracy (method of integration & step size) and precision (double or float)

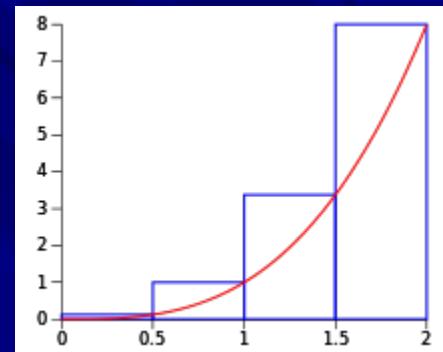
# Methods of Integration (Riemann Sum)

Left Riemann Sum  
**(Forward)**  
 $A_n = f(a) dx$   
 $dx = [b-a]/n$



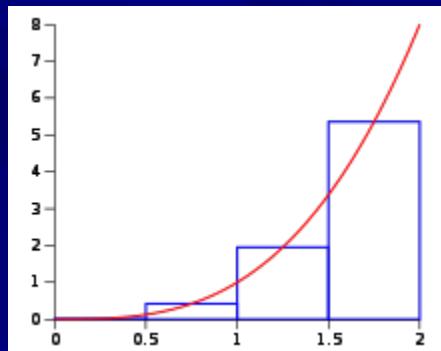
$$A_{b-a} = dx [ (f(a) + f(a+dx) + f(a+2dx) + \dots + f(b-dx)) ]$$

Right Riemann Sum  
**(Backward)**  
 $A_n = f(a+dx) dx$   
 $dx = [b-a]/n$



$$A_{b-a} = dx [ f(a+dx) + f(a+2dx) + \dots + f(b) ]$$

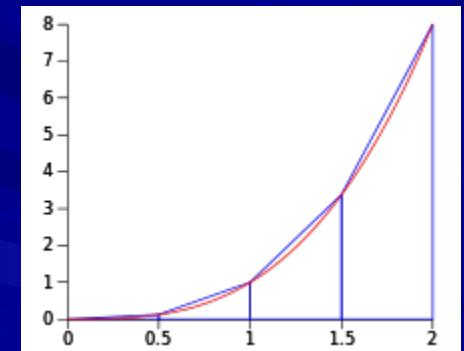
Midpoint Rule  
 $A_n = f(a+dx/2) dx$   
 $dx = [b-a]/n$



$$A_{b-a} = dx [ f(a+dx/2) + f(a+3dx/2) + \dots + f(b-dx/2) ]$$

Trapezoidal Rule

$$A_{\text{trap}} = \frac{1}{2} h (f(a)+f(b))$$
$$dx = [b-a]/n$$
$$A_n = \frac{1}{2} dx (f(a)+f(a+dx))$$



$$\frac{1}{2} dx [ (f(a) + f(a+dx)) + (f(a+dx) + f(a+2dx)) \dots ]$$
$$\frac{1}{2} dx [ f(a) + 2f(a+dx) + 2f(a+2dx) + \dots + f(b) ]$$

Starter code: [https://www.ecst.csuchico.edu/~sbsiewert/csci551/code/hello\\_integrators/](https://www.ecst.csuchico.edu/~sbsiewert/csci551/code/hello_integrators/)

# How do I verify Numerical Integration?

- Test with definite integral from Calculus – exact solution
- E.g., area under  $\sin(x)$  over well-defined interval (2.0)

Known anti-derivative of sine and evaluation

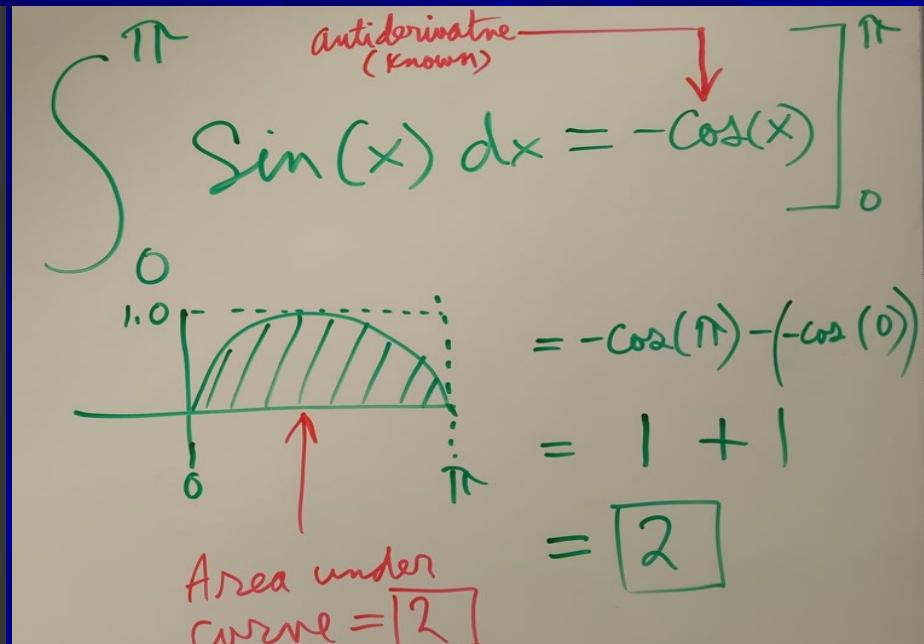
- Definite integral solution
- We know from calculus
- Can be proven

Use this fact to verify your numerical integration

- Riemann Sum (function)
- Trapezoidal (function)
- Simpson's Rule (function)
- Runge-Kutta (differential equations)
- Other methods of quadrature (math topic)
- Estimation of area under a curve (function)
- Integration in general of a differential equation
- In applications, the functions generally represent physical quantities, the derivatives represent their rates of change, and the differential equation defines a relationship between the two.

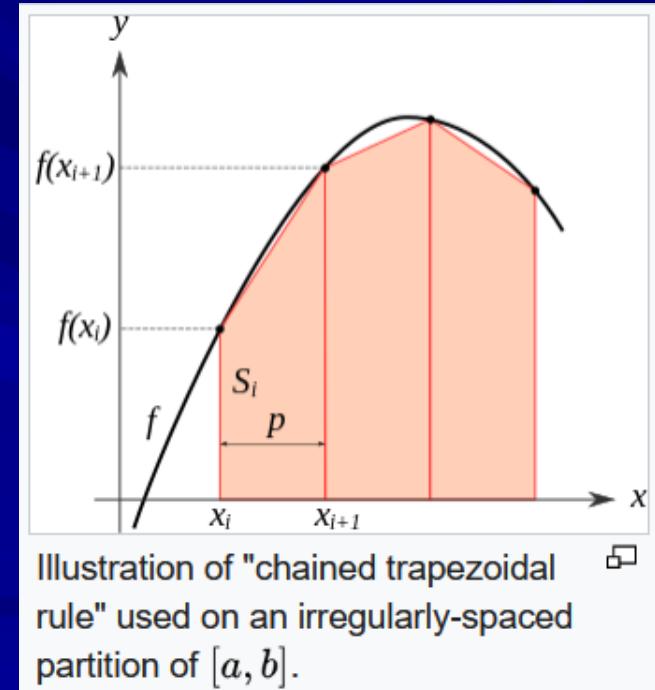
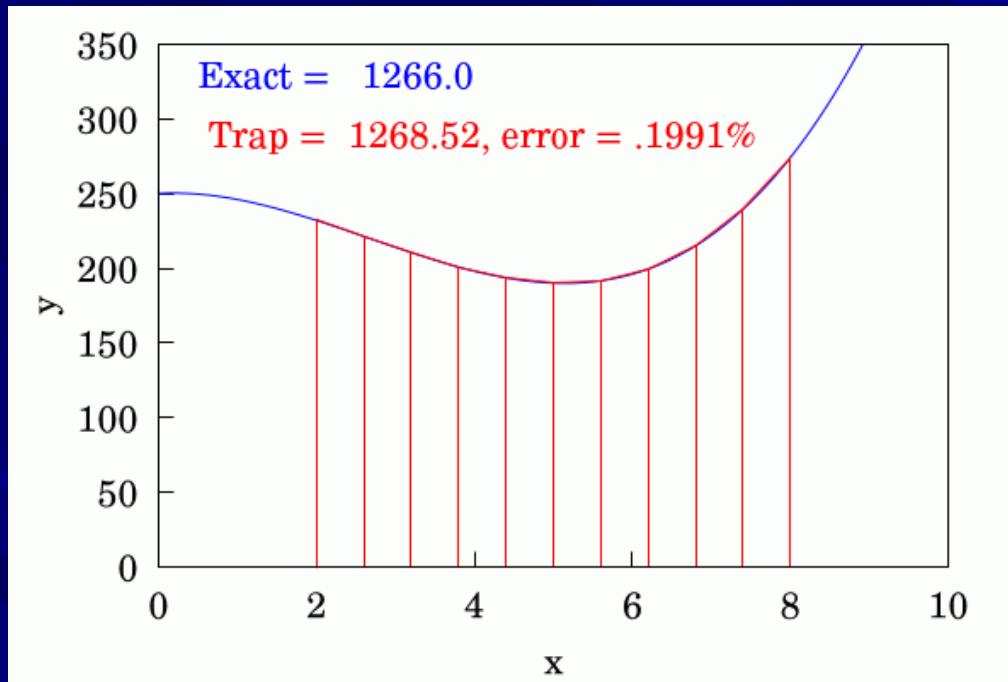
Unit Circle integration of velocity of shadow

- Assume light source on vertical axis
- Velocity of the shadow is  $\sin(\theta)$
- Position is the Integral of velocity
- Distance is 2.0 for unit circle



# Trapezoidal Rule

- Consider following for error in numerical integration
  - Step size (smaller can reduce error)
  - Method of estimating area over interval (can be more accurate)
  - Difficult to prove one method is always best – depends on function, increases complexity of code, could be fixed or variable



[https://en.wikipedia.org/wiki/Trapezoidal\\_rule](https://en.wikipedia.org/wiki/Trapezoidal_rule)

# More Advanced Methods

## Trapezoidal Rule (functions)

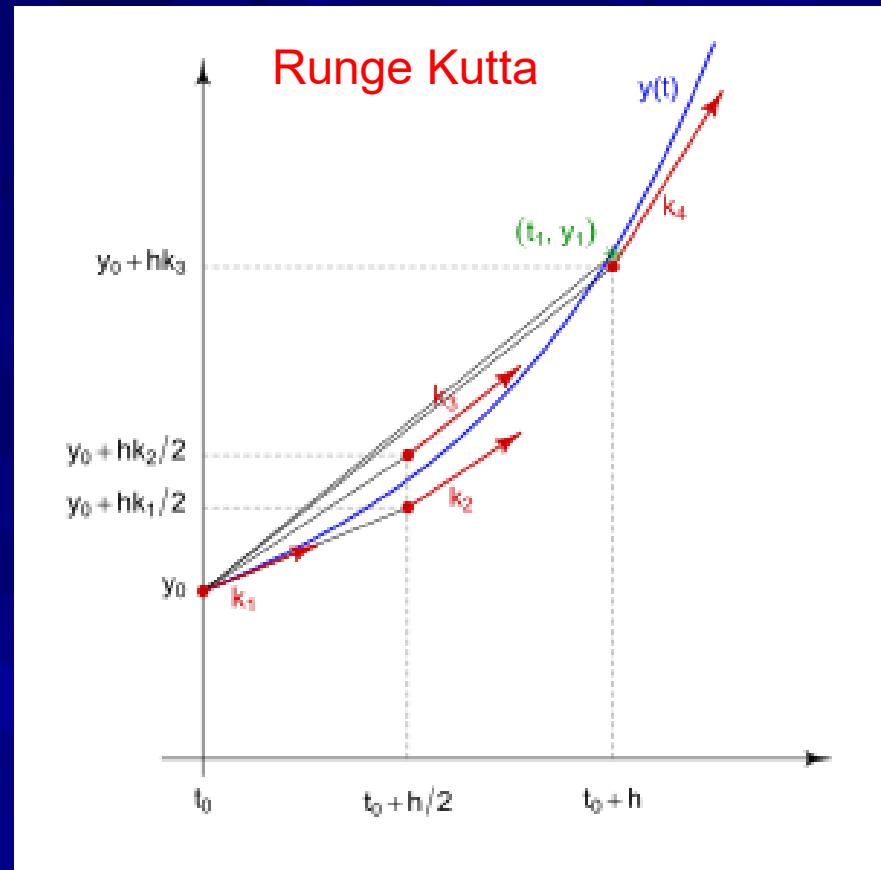
- Can be better than Riemann for same step size

## Simpson's Rule (functions)

- Can be better than trapezoidal for same step size

## Runge-Kutta methods (differential equations)

- Fixed step size, 4 increments for RK4
- We are estimating 4 slopes over the interval
- We are averaging  $k_1$ ,  $k_2$ ,  $k_3$  and  $k_4$  and weighting the  $k_2$  and  $k_3$  by 2x



$$y_{n+1} = y_n + \frac{1}{6}h(k_1 + 2k_2 + 2k_3 + k_4),$$
$$t_{n+1} = t_n + h$$

[https://en.wikipedia.org/wiki/Runge%20Kutta\\_methods](https://en.wikipedia.org/wiki/Runge%20Kutta_methods)

# Summary – Numerical Integration

- Use Simple Riemann Sums – Left, Right, Mid-point
- Consider Trapezoidal and Simpson's
- Step size matters!
  - Smaller is generally better, but requires more computation
  - Think of step size as you do “resolution” for images
- Advanced methods might allow for larger step sizes or even variable step sizes
- Differential and difference equations ([for the curious](#)) are interesting for more advanced
  - e.g., [Lorenz Equations](#), [N-body simulation](#)
  - Has  $y$  and  $dy/dt$  in same equation (can have  $d^2y/dt^2$  for second order)
- **Simple integration useful for any problem that involves the area under a curve** (function), e.g., our train problems
- Equations of motion – can be simple integration problem (e.g., our 1-DOF train problem), or more complex differential equations depending on system (e.g., [spring mass damper system](#))